

# Prediction for Average Bottom hole Pressure in Volve field using Lasso Regression Model



SAVYSTAT

Notebook    TypeError

...

```
import pandas as pd
import numpy as np
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Lasso
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
```

```
df = pd.read_excel("Volve field dataset.xlsx")
```

```
df.head()
for column_name in df.columns:
    print(column_name)
```

```
ON_STREAM_HRS
AVG_DOWNHOLE_TEMPERATURE
AVG_ANNULUS_PRESS
AVG_CHOKE_SIZE_P
AVG_WHP_P
AVG_WHT_P
DP_CHOKE_SIZE
BORE_OIL_VOL
BORE_GAS_VOL
AVG_DOWNHOLE_PRESSURE
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15634 entries, 0 to 15633
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ON_STREAM_HRS                        15349 non-null  float64
1   AVG_DOWNHOLE_TEMPERATURE             8980 non-null  float64
2   AVG_ANNULUS_PRESS                    7890 non-null  float64
3   AVG_CHOKE_SIZE_P                     8919 non-null  float64
4   AVG_WHP_P                            9155 non-null  float64
5   AVG_WHT_P                            9146 non-null  float64
6   DP_CHOKE_SIZE                        15340 non-null  float64
7   BORE_OIL_VOL                         9161 non-null  float64
8   BORE_GAS_VOL                         9161 non-null  float64
9   AVG_DOWNHOLE_PRESSURE                8980 non-null  float64
dtypes: float64(10)
memory usage: 1.2 MB
```

```
# Check for missing values
missing_values = df.isnull().sum()
print(missing_values)
```

```
ON_STREAM_HRS                285
AVG_DOWNHOLE_TEMPERATURE     6654
AVG_ANNULUS_PRESS            7744
AVG_CHOKE_SIZE_P             6715
AVG_WHP_P                    6479
AVG_WHT_P                    6488
DP_CHOKE_SIZE                 294
BORE_OIL_VOL                  6473
BORE_GAS_VOL                  6473
AVG_DOWNHOLE_PRESSURE        6654
dtype: int64
```

```
# Drop rows containing zeros
df = df[(df != 0).all(axis=1)]
# # Remove rows with any missing values
df_cleaned = df.dropna()
df_cleaned
```

	ON_STREAM_HRS	AVG_DOWNHOLE_TEMPERATURE	AVG_ANNULUS_PRESS	AVG_CHOKE_SIZE_P	
762	7.00000	0.352200	2.885360	3.256548	
763	24.00000	60.315740	19.464510	8.549131	1
769	5.07514	105.551370	21.550252	2.540804	
772	15.07486	104.933215	1.652926	6.116182	
773	24.00000	105.439765	17.308850	9.951288	
...	...	...	...	...	
8923	24.00000	106.517574	21.318431	31.575767	
8924	24.00000	106.515586	21.105330	31.540612	
8925	24.00000	106.521356	21.353661	31.522096	
8926	24.00000	106.506781	20.629658	31.523457	
8927	18.37500	106.507232	20.404848	24.922565	

4166 rows × 10 columns

Next steps:

Generate code with df\_cleaned

 View recommended plots

```
missing_values = df_cleaned.isnull().sum()
print(missing_values)
```

```
ON_STREAM_HRS      0
AVG_DOWNHOLE_TEMPERATURE  0
AVG_ANNULUS_PRESS  0
AVG_CHOKE_SIZE_P    0
AVG_WHP_P           0
AVG_WHT_P           0
DP_CHOKE_SIZE       0
BORE_OIL_VOL        0
BORE_GAS_VOL        0
AVG_DOWNHOLE_PRESSURE 0
dtype: int64
```

```
df = df_cleaned
```

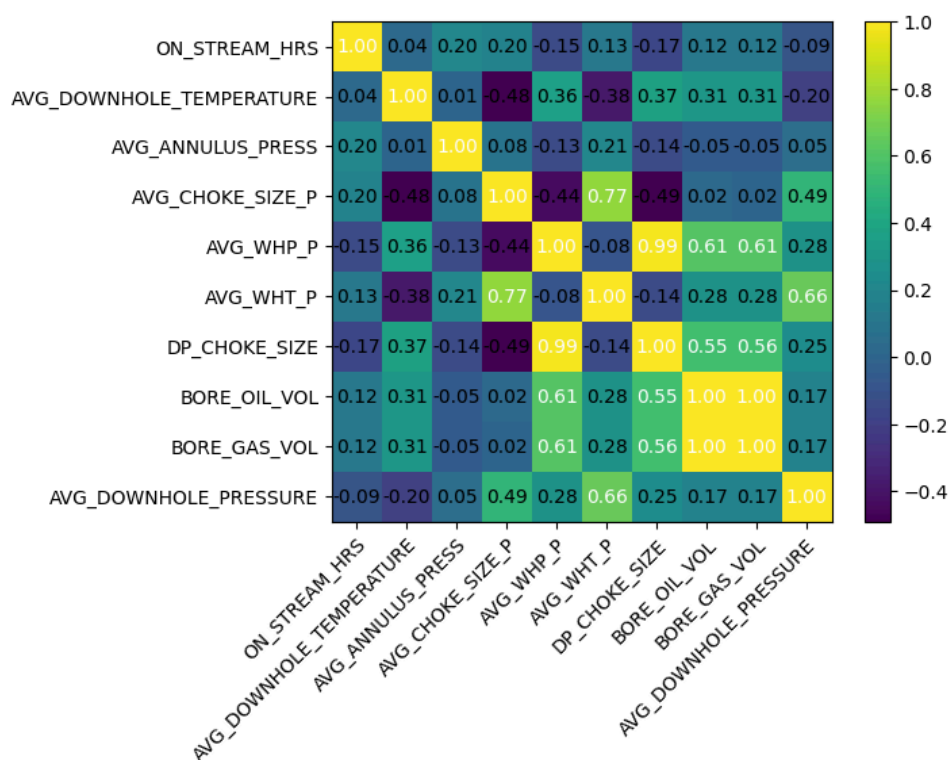
```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 4166 entries, 762 to 8927
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ON_STREAM_HRS                        4166 non-null   float64
1   AVG_DOWNHOLE_TEMPERATURE             4166 non-null   float64
2   AVG_ANNULUS_PRESS                   4166 non-null   float64
3   AVG_CHOKE_SIZE_P                    4166 non-null   float64
4   AVG_WHP_P                           4166 non-null   float64
5   AVG_WHT_P                           4166 non-null   float64
6   DP_CHOKE_SIZE                       4166 non-null   float64
7   BORE_OIL_VOL                        4166 non-null   float64
8   BORE_GAS_VOL                        4166 non-null   float64
9   AVG_DOWNHOLE_PRESSURE               4166 non-null   float64
dtypes: float64(10)
memory usage: 358.0 KB
```

```
df.describe()
```

	ON_STREAM_HRS	AVG_DOWNHOLE_TEMPERATURE	AVG_ANNULUS_PRESS	AVG_CHOKE_SIZE_P
<b>count</b>	4166.000000	4166.000000	4166.000000	4166.000000
<b>mean</b>	23.176426	104.105784	18.317129	53.508996
<b>std</b>	3.227762	4.097660	5.121428	37.138717
<b>min</b>	0.250000	0.352200	0.000020	0.600000
<b>25%</b>	24.000000	100.475027	14.382808	12.310157
<b>50%</b>	24.000000	105.944562	18.879005	50.840645
<b>75%</b>	24.000000	106.481651	22.065755	100.000000
<b>max</b>	25.000000	107.507552	30.019828	100.000000

```
import matplotlib.pyplot as plt
from mlxtend.plotting import heatmap
cols = df.columns.tolist()
cm = np.corrcoef(df[cols].values.T)
hm = heatmap(cm, row_names = cols, column_names = cols)
# Enlarge the figure size
plt.figure(figsize=(25, 20))
plt.show()
```



<Figure size 2500x2000 with 0 Axes>

```
x = df.iloc[:, :-1]
y = df.iloc[:, -1]

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, random_state=0)

x_train.shape

(3124, 9)

y_train.shape
```

```

(3124,)

import numpy as np

# Take the natural logarithm of the features
x_train_log = np.log(x_train)
x_test_log = np.log(x_test)

import numpy as np
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Lasso

# Generate a range of alpha values spanning multiple orders of magnitude
alphas = np.logspace(-3, 3, 7) # This will generate values from 0.001 to 1000

# Create a pipeline with feature scaling and Lasso regression
pipeline = Pipeline([
    ('scaler', StandardScaler()), # Feature scaling
    ('lasso', Lasso()) # Lasso regression
])

# Define hyperparameters to tune
param_grid = {'lasso__alpha': alphas}

# Perform grid search with 5-fold cross-validation
grid_search = GridSearchCV(estimator=pipeline, param_grid=param_grid, cv=5, scoring='neg_mean_squared_error')

# Fit grid search to the training data
grid_search.fit(x_train_log, y_train) # Replace x_train_scaled and y_train with your training data

# Get the best alpha value
best_alpha = grid_search.best_params_['lasso__alpha']

# Print the best alpha value
print("Best alpha:", best_alpha)

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning: Objective
model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning: Objective
model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning: Objective
model = cd_fast.enet_coordinate_descent(
Best alpha: 0.1

#fit the model
from sklearn.linear_model import Lasso
lasso_mod = Lasso(alpha=0.1).fit(x_train_log,y_train)

# lasso_mod = Lasso(alpha=0.01, max_iter=10000).fit(x_train_log, y_train)

lasso_mod.coef_

array([ 2.03044613, 30.39507044, -0.33456594, 16.662728 ,
        50.75392338, 38.25917017,  3.93565611, -18.32982729,
        -0.          ])

print("Intercept:", lasso_mod.intercept_)


Intercept: -206.58070093289825

# identifying which of the features contribute to the model
# Get feature names
feature_names = df.columns
coefficients = lasso_mod.coef_

# Create DataFrame to store coefficients and feature names
coefficients_df = pd.DataFrame({'Feature': feature_names[:-1], 'Coefficient': coefficients})

```

```
# Plot coefficients
plt.figure(figsize=(12, 8))
plt.bar(coefficients_df['Feature'], coefficients_df['Coefficient'])
plt.xlabel('Features')
plt.ylabel('Coefficient Value')
plt.title('Coefficients at Alpha = 0.01')
plt.xticks(rotation=90) # Rotate x-axis labels for better readability
plt.grid(True)
plt.show()
```



```
#printing the r-squared on training data
print('R- squared on training data', lasso_mod.score(x_train_log, y_train)*100)
```

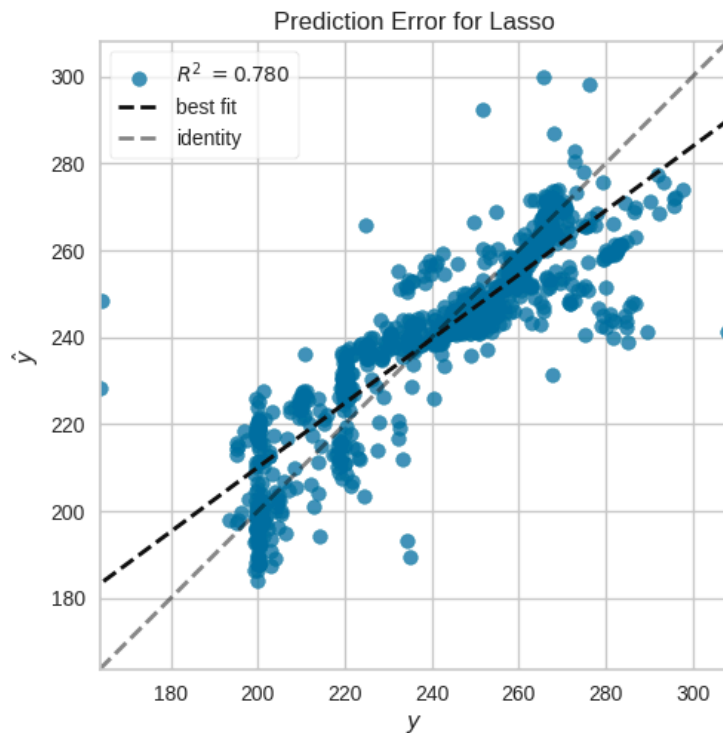
R- squared on training data 76.58081637841275

```
#printing the r-squared on test data
print('R- squared on test data', lasso_mod.score(x_test_log, y_test)*100)
```

R- squared on test data 78.00258393658079

```
#plotting the prediction errors and residuals using yellow brick
!pip install yellowbrick
from sklearn.preprocessing import StandardScaler
from yellowbrick.regressor import PredictionError, ResidualsPlot
visualizer = PredictionError(lasso_mod)
visualizer.fit(x_train_log, y_train)
visualizer.score(x_test_log, y_test)
visualizer.poof()
```


Requirement already satisfied: yellowbrick in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: matplotlib!=3.0.0,>=2.0.2 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: scikit-learn>=1.0.0 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: cycycler>=0.10.0 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages  
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (f  
 /usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not  
 warnings.warn()



```
<Axes: title={'center': 'Prediction Error for Lasso'}, xlabel='$y$',  
ylabel='$\hat{y}$'>
```

```
#plotting the residuals
visualizer = ResidualsPlot(lasso_mod)
visualizer.fit(x_train_log, y_train)
visualizer.score(x_test_log, y_test)
visualizer.poof()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not
warnings.warn(
```



```
<Axes: title={'center': 'Residuals for Lasso Model'}, xlabel='Predicted Value',
ylabel='Residuals'>
```

```
y_pred =lasso_mod.predict(x_test_log)
```

```
# Create a DataFrame to display the predicted values for the test data
```

```
result_df = pd.DataFrame({'Predicted Values': y_pred})
```

```
# Create a DataFrame to display the predicted and actual values
```

```
result_df = pd.DataFrame({'Actual Values': y_test, 'Predicted Values of AVG_DOWNHOLE_PRESSURE': y_pred})
```

```
# Display the DataFrame
```

```
print(result_df.head(15))
```

	Actual Values	Predicted Values of AVG_DOWNHOLE_PRESSURE
1163	265.812269	245.906340
8087	199.714866	189.428519
7184	265.382097	262.564390
2170	235.708089	252.861815
8683	223.130863	212.361288
2859	267.022065	253.122667
1860	251.343919	241.469494
7640	267.210376	267.793313
7527	266.157293	261.145763
7397	262.291793	251.654462
7773	267.212864	266.849270
1239	219.749815	226.555836
8072	200.341786	187.681363
1247	219.458958	228.378539
1079	225.209442	233.826684

```
#For the test data
```

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_absolute_percentage_error,r2_score
```

```
y_pred = lasso_mod.predict(x_test_log)
```

```
# Calculate MAE
```

```
mae_test = mean_absolute_error(y_test, y_pred)
```

```
# Calculate MSE
```

```
mse_test = mean_squared_error(y_test, y_pred)
```

```
# Calculate RMSE
```

```
rmse_test = np.sqrt(mse_test)
```



```
#calculate MAPE
mape_test = mean_absolute_percentage_error(y_test,y_pred)
#calculate R squared
rsq = r2_score(y_test, y_pred)
# Print MAE, MSE, RMSE, R squared
print("Mean Absolute Error (MAE): ", mae_test)
print("Mean Squared Error (MSE): ", mse_test)
print("Root Mean Squared Error (RMSE): ", rmse_test)
print("Mean absolute percentage error: ", mape_test)
print("R squared: ", rsq)
```

```
Mean Absolute Error (MAE): 8.255949208554767
Mean Squared Error (MSE): 139.6476252955515
Root Mean Squared Error (RMSE): 11.817259635615674
Mean absolute percentage error: 0.034813327678556105
R squared: 0.7800258393658079
```

```
!pip install nbconvert
```

```
Requirement already satisfied: nbconvert in /usr/local/lib/python3.10/dist-packages (6.5.4)
Requirement already satisfied: lxml in /usr/local/lib/python3.10/dist-packages (from nbconvert) (4.9.4)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (4.12.3)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from nbconvert) (6.1.0)
Requirement already satisfied: defusedxml in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.7.1)
Requirement already satisfied: entrypoints>=0.2.2 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.4)
Requirement already satisfied: Jinja2>=3.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (3.1.4)
Requirement already satisfied: jupyter-core>=4.7 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (5.7.2)
Requirement already satisfied: jupyterlab-pygments in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (2.1.5)
Requirement already satisfied: mistune<2,>=0.8.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.8.4)
Requirement already satisfied: nbclient>=0.5.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.10.0)
Requirement already satisfied: nbformat>=5.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (5.10.4)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from nbconvert) (24.0)
Requirement already satisfied: pandocfilters>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (1.5.1)
Requirement already satisfied: pygments>=2.4.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (2.16.1)
Requirement already satisfied: tinycss2 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (1.3.0)
Requirement already satisfied: traitlets>=5.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (5.7.1)
Requirement already satisfied: platformdirs>=2.5 in /usr/local/lib/python3.10/dist-packages (from jupyter-core>=4.7->nbconvert) (4.2.2)
Requirement already satisfied: jupyter-client>=6.1.12 in /usr/local/lib/python3.10/dist-packages (from nbclient>=0.5.0->nbconvert) (7.2.0)
Requirement already satisfied: fastjsonschema>=2.15 in /usr/local/lib/python3.10/dist-packages (from nbformat>=5.1->nbconvert) (2.19.1)
Requirement already satisfied: jsonschema>=2.6 in /usr/local/lib/python3.10/dist-packages (from nbformat>=5.1->nbconvert) (4.19.0)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4->nbconvert) (2.5)
```